
CNN-based Ship Localization in Satellite Imagery: A U-Net Approach for the Airbus Ship Detection Challenge

Palásti András

Student at Budapest University of Technology and Economics
andras.palasti@edu.bme.hu

Kurcsi Norbert

Student at Budapest University of Technology and Economics
kurcsi.norbert@edu.bme.hu

Wittmajer Dávid

Student at Budapest University of Technology and Economics
wittmajerd@edu.bme.hu

Abstract

In the context of the Airbus Ship Detection Challenge, searching for ships on the seas is akin to finding a needle in a haystack. Our paper addresses this challenge by using a U-Net image segmentation model that relies solely on convolutional layers to rapidly and accurately localize ships in satellite images. This approach holds promising potential for enhancing ship detection technology and advancing the field of satellite image analysis.

1 Introduction

We developed our model for the Airbus Ship Detection Challenge, tasked with swiftly and accurately identifying ships in satellite imagery. The objective is to facilitate enhanced monitoring of oceans amidst increasing maritime traffic, allowing environmental organizations, insurers, and governmental authorities to better oversee potential incidents such as accidents, pirate attacks, illicit drug and cargo transportation, and illegal fishing. Airbus provides crucial maritime observation services, emphasizing the need for precise and prompt information delivery to enhance the efficiency of their clients.

The primary challenge lies in detecting ships that are obscured by clouds, navigating adverse weather conditions, or are exceptionally small. Moreover, identifying vessels within harbors amid diverse structures and terrain poses an additional difficulty. This paper outlines our innovative approach to address these challenges and contribute to advancing maritime surveillance capabilities.

2 Related work

Convolutional Neural Networks (CNNs) in Object Detection Recent advancements in object detection have been significantly influenced by Convolutional Neural Networks (CNNs). Ren et al. (2015) introduced Faster R-CNN, a seminal work that integrates region proposal networks with CNNs for efficient object detection, laying the foundation for subsequent improvements in accuracy and speed [1].

Semantic Segmentation with U-Net The U-Net architecture, proposed by Ronneberger et al. (2015), initially gained prominence in the field of medical image segmentation. Its application has since expanded due to its efficacy and simplicity. U-Net’s encoder-decoder structure, characterized by contracting and expanding pathways, has been widely adopted in diverse segmentation tasks [2].

Fully Convolutional Networks (FCN) The work of Long, Shelhamer, and Darrell (2014) introduced Fully Convolutional Networks (FCN) for semantic segmentation, marking a significant milestone in end-to-end pixelwise labeling. FCN has laid the foundation for subsequent developments in semantic segmentation, influencing the design of numerous segmentation architectures [3].

Instance Segmentation Evaluation The work of Dollar et al. (2014) on the COCO (Common Objects in Context) dataset has served as a benchmark for evaluating instance segmentation algorithms. The dataset provides a diverse set of images with detailed annotations, facilitating rigorous evaluation of segmentation performance [4].

ImageNet Pre-trained Models Razavian et al. (2014) explored the effectiveness of leveraging pre-trained CNN models on ImageNet for image classification tasks. The use of pre-trained models has shown promise in feature extraction for various applications, including remote sensing and satellite image analysis [5].

These diverse approaches collectively contribute to the growing body of literature in image segmentation, providing insights and methodologies that address different challenges in the field.

3 Architecture




Figure 1: The original UNet architecture was employed, albeit in a modified and simplified version. In this adaptation, the width and height of the input image tile align with the corresponding dimensions of the output segmentation map.

In our approach to the task, we employed a slightly modified UNet architecture, as depicted in Figure 1. The sole deviation in our implementation lies in the absence of padding on the input image tile. Consequently, the input size aligns with the dimensions of the output segmentation, specifically matching in width and height. This was achieved by replacing the ‘copy and crop’ operations in the original model. Instead, we first padded the smaller image with zeros to match the size of the larger image and subsequently concatenated the two.

4 Implementation

4.1 Data preparation

Originally, the Airbus Ship Detection dataset comprised 192556 images. Out of these, only 42556 images contained ships, accounting for less than 25% of the dataset. Given the enormity of the full dataset for our specific requirements, and considering that the majority of images did not feature ships, we opted to utilize a subset of 60000 images. This subset included all images containing ships (42556 images) and additional images without ships ($60000 - 42556 = 17444$ images).

We divided these images into training, validation, and test sets with ratios of 90%, 5%, and 5%, respectively. Our dataset is publicly available at the following URL: <https://drive.google.com/uc?id=1V-oxFZhctefBXL4noEgNG-ENaHoszTJE>.

4.2 Training




Figure 2: The top row depicts ground truth segmentations, and the bottom row shows predicted masks. The left side of the images corresponds to segmentations produced by a network trained on full-sized images, while the right side displays predictions from a network trained on cropped images.

For our initial solution, we employed the UNet model as described without engaging in intricate procedures. The model was trained using provided training images and corresponding segmentation masks, with input images sized 768×768 . The outcomes of this model are illustrated in Figure 2.

Our conjecture regarding the model’s performance is that, when trained on full-sized images, it struggles to discern crucial features due to the expansive context.

To address this issue, during training, we adopted a strategy of cropping a 256×256 region from the original training set image. If the image contains a ship, that region retains the ship; otherwise, a random portion of the image is cropped. This approach ensures that our model receives training data in a more focused manner.

4.3 Evaluation

For our network’s evaluation strategy, we considered two options: either performing predictions on the full-sized image or splitting the image into 256×256 parts and conducting predictions on each part. We experimented with both approaches, and the results are presented in Table 1.

As evident from Table 1, it may appear that our optimal strategy involves dividing the image into 256×256 regions and conducting inference on each region. However, this observation is misleading

Table 1: Results

Inference on	IoU Score	Dice Score
Full-sized images	63.34%	72.88%
256×256 regions	86.09%	87.90%
256×256 regions that contain ships	60.68%	69.39%
Full-sized images with model trained on full sized images	54.54%	63.39%

because our evaluation includes images that do not contain any ships. Many of these ship-absent images achieve a perfect score of 1.0, thereby inflating the mean of the scores. This is further substantiated by the fact that our results exhibit a significant decline when we exclusively perform predictions on images containing ships.

5 Conclusion

In conclusion, our research highlights the crucial impact of training strategies on the performance of image segmentation models. Through experimentation with different training methodologies, we observed a significant improvement in results when the model was exposed to the most relevant portions of the data.

This approach proved to be particularly effective, leading to enhanced segmentation performance. The emphasis on training the model with images containing ships, and carefully adjusting the segmentation masks accordingly, resulted in superior inference outcomes.

6 Future improvements

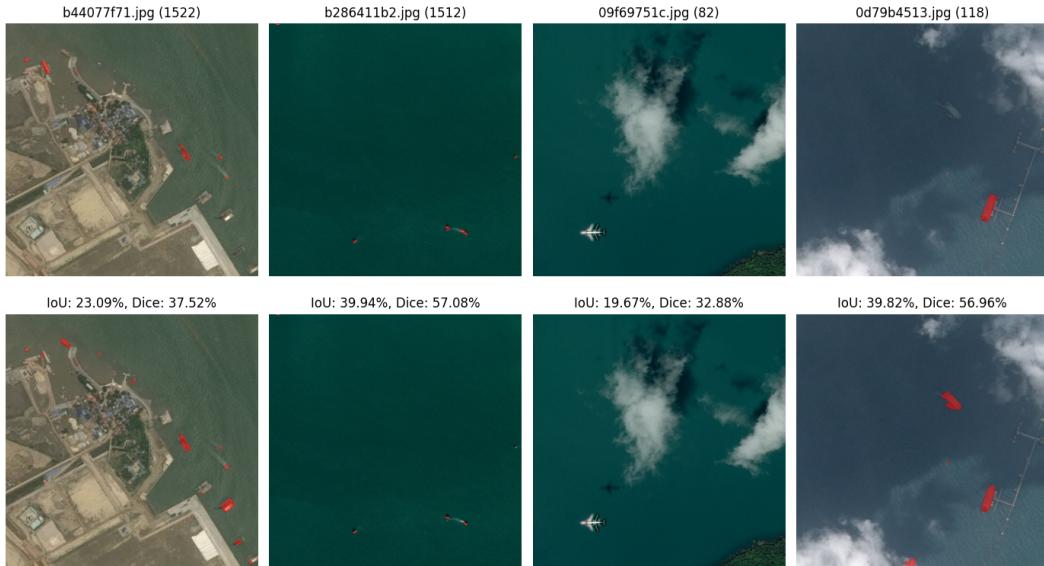


Figure 3: Some of the images where our model fails to achieve respectable results.

While our current model demonstrates promising results, there are areas where improvements could significantly enhance its performance. A careful examination of the worst segmentations, as depicted in Figure 3, reveals challenges in scenarios where land is present. Additionally, the model tends to struggle in distinguishing between harbors and ships.

To address these weaknesses, a potential avenue for improvement involves a two-step training approach. Initially, training on cropped images could allow the model to focus on learning the

intricate details relevant to ship detection. Subsequently, fine-tuning the model on full-sized images may help it generalize better to diverse scenes, including those with land and harbor complexities.

This sequential training strategy aims to leverage the advantages of both cropped and full-sized images, allowing the model to capture intricate details while maintaining a broader understanding of the overall context.

References

- [1] Ren, S., He, K., Girshick, R., & Sun, J. (2015). Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks. In Advances in neural information processing systems (pp. 91-99).
- [2] Ronneberger, O., Fischer, P., & Brox, T. (2015). U-Net: Convolutional Networks for Biomedical Image Segmentation. In International Conference on Medical image computing and computer-assisted intervention (pp. 234-241).
- [3] Long, J., Shelhamer, E., & Darrell, T. (2015). Fully Convolutional Networks for Semantic Segmentation. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (pp. 3431-3440).
- [4] Dollar, P., Wojek, C., Schiele, B., & Perona, P. (2012). Pedestrian Detection: An Evaluation of the State of the Art. *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, 34(4), 743-761.
- [5] Razavian, A. S., Azizpour, H., Sullivan, J., & Carlsson, S. (2014). CNN Features off-the-shelf: an Astounding Baseline for Recognition. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops (pp. 806-813).